

# Next Word Prediction Using Stacked LSTM Networks in Deep Learning

The text dataset is a children stories which is taken from kaggle. Using this dataset, a next word predicting model has been trained using stacked LSTM.

```
In [3]: !kaggle datasets download -d edenbd/children-stories-text-corpus
```

Dataset URL: <https://www.kaggle.com/datasets/edenbd/children-stories-text-corpus>

License(s): CC0-1.0

children-stories-text-corpus.zip: Skipping, found more recently modified local copy (use --force to force download)

```
In [4]: import zipfile
zip_ref=zipfile.ZipFile("children-stories-text-corpus.zip","r")
zip_ref.extractall()
zip_ref.close()
```

```
In [5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import re
```

```
In [6]: with open('cleaned_merged_fairy_tales_without_eos.txt', 'r', encoding='UTF-8') as file:
        content = file.read()
        print(content[:500])
```

The Happy Prince.

HIGH above the city, on a tall column, stood the statue of the Happy Prince. He was gilded all over with thin leaves of fine gold, for eyes he had two bright sapphires, and a large red ruby glowed on his sword-hilt.

He was very much admired indeed. "He is as beautiful as a weathercock," remarked one of the Town Councillors who wished to gain a reputation for having artistic tastes; "only not quite so useful," he added, fearing lest people should think him impractical, which h

```
In [7]: #removing some punctuations using regex
def reg(text):
    pattern=r"^[^w\s\.\,\;]"
    processed=re.sub(pattern,'',text)
    return processed

reg_text=reg(content)
```

This is a huge datasets for a next word predicting model. First 1000000 strings have taken from it so that it can be handled by the system and also to train it quickly.

```
In [9]: reg_text=reg_text.replace("\n"," ")
txt=reg_text[:1000000]
txt[:1000]
```

```
Out[9]: 'The Happy Prince. HIGH above the city, on a tall column, stood the statue of the Happy Prince. He was gilded all over with thin leaves of fine gold, for eyes he had two bright sapphires, and a large red ruby glowed on his swordhilt. He was very much admired indeed. He is as beautiful as a weathercock, remarked one of the Town Councillors who wished to gain a reputation for having artistic tastes; only not quite so useful, he added, fearing lest people should think him impractical, which he really was not. Why cant you be like the Happy Prince asked a sensible mother of her little boy who was crying for the moon. The Happy Prince never dreams of crying for anything. I am glad there is some one in the world who is quite happy, muttered a disappointed man as he gazed at the wonderful statue. He looks just like an angel, said the Charity Children as they came out of the cathedral in their bright scarlet cloaks and their clean white pinafores. How do you know said the Mathematical Master'
```

```
In [10]: len(txt)
```

```
Out[10]: 1000000
```

```
In [11]: import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.utils import to_categorical
```

For vectorization, "Tokenizer" object has been used. "fit\_on\_texts" function allocated a number for each word in that text.

```
In [13]: token=Tokenizer()
token.fit_on_texts([txt])
```

```
In [80]: # Showing index of word
token_word=token.word_index
```

"texts\_to\_sequences" function converted the text into a list of tokens corresponding the word.

```
In [16]: token_line=token.texts_to_sequences([txt])[0]
token_len=len(token_line)
token_len
```

```
Out[16]: 191738
```

```
In [82]: token_line[:15]
```

```
Out[82]: [1, 209, 72, 302, 444, 1, 583, 29, 4, 1030, 3394, 182, 1, 1283, 5]
```

The whole token sequences have been made a array having (n,4) dimension. So that for every row first three elements can be used as input and fourth element as output.

```
In [86]: input=[]
for i in range(4*(token_len//4)-3):
    input.append(token_line[i:i+4])
input[:10]
```

```
Out[86]: [[1, 209, 72, 302],
[209, 72, 302, 444],
[72, 302, 444, 1],
[302, 444, 1, 583],
[444, 1, 583, 29],
[1, 583, 29, 4],
[583, 29, 4, 1030],
[29, 4, 1030, 3394],
[4, 1030, 3394, 182],
[1030, 3394, 182, 1]]
```

Here the splitting has been done which was stated above. "x" is the input and "y\_" is the output.

```
In [21]: x=[]
y_=[]

for i in input:
    x.append(i[:-1])
    y_.append(i[-1])
x=np.array(x)
y_=np.array(y_)
print(f"shape of x: {x.shape}")
print(f"shape of y: {y_.shape}")
```

```
shape of x: (191733, 3)
shape of y: (191733,)
```

```
In [22]: x_col=x.shape[1]
print(f"total column in x: {x_col}")
```

```
total column in x: 3
```

We can notice that token started from 1. So to include 0, "+1" business has been done here.

```
In [24]: total_word=len(token.word_index)
total_word=total_word+1
total_word
```

```
Out[24]: 9234
```

This is a multiclass classification. So the output can't be just a number. So, the categorization has been done using "to\_categorical" function. Also notice that, the output can be any word from our word vocabulary. so, the "total\_word" has been taken as length of one output.

```
In [26]: y=to_categorical(y_,num_classes=total_word)
y.shape
```

```
Out[26]: (191733, 9234)
```

```
In [27]: from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, LSTM, Embedding,Dropout
```

1. Word embedding has been done to create a array corresponding every unique word/token.
2. Then LSTM layers have been used for finding the pattern.
3. To reduce overfitting dropout has been used.
4. As activation function, Softmax has been used because it is a multiclass classification problem.

```
In [29]: model=Sequential()

model.add(keras.Input(shape=(x_col,)))
model.add(Embedding(total_word,100))
model.add(LSTM(500, return_sequences=True))
model.add(LSTM(500))
model.add(Dropout(0.5))
model.add(Dense(total_word,activation="softmax"))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 3, 100)	923,400
lstm (LSTM)	(None, 3, 500)	1,202,000
lstm_1 (LSTM)	(None, 500)	2,002,000
dropout (Dropout)	(None, 500)	0
dense (Dense)	(None, 9234)	4,626,234

**Total params:** 8,753,634 (33.39 MB)

**Trainable params:** 8,753,634 (33.39 MB)

**Non-trainable params:** 0 (0.00 B)

```
In [30]: model.compile(optimizer="adam",loss="categorical_crossentropy",metrics=["accuracy"])
```

```
In [31]: # Using earlystopping to save from overfitting
```


```
callback=keras.callbacks.EarlyStopping(
    monitor="loss",
    min_delta=0,
    patience=0,
    verbose=1,
    mode="auto",
    baseline=None,
)
```

```
In [32]: history=model.fit(x,y,epochs=100, batch_size=50, callbacks=callback) #, validation_data=(x_test,y_test)
```


Epoch 1/100			
3835/3835	<div></div>	91s	23ms/step - accuracy: 0.0769 - loss: 6.3871
Epoch 2/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.1233 - loss: 5.5513
Epoch 3/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.1438 - loss: 5.2408
Epoch 4/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.1580 - loss: 5.0369
Epoch 5/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.1695 - loss: 4.8613
Epoch 6/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.1847 - loss: 4.6762
Epoch 7/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.1975 - loss: 4.5126
Epoch 8/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.2139 - loss: 4.3399
Epoch 9/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.2347 - loss: 4.1209
Epoch 10/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.2541 - loss: 3.9301
Epoch 11/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.2750 - loss: 3.7294
Epoch 12/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.2982 - loss: 3.5363
Epoch 13/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.3224 - loss: 3.3371
Epoch 14/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.3485 - loss: 3.1553
Epoch 15/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.3698 - loss: 2.9869
Epoch 16/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.3923 - loss: 2.8340
Epoch 17/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.4157 - loss: 2.6872
Epoch 18/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.4401 - loss: 2.5372
Epoch 19/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.4586 - loss: 2.4299
Epoch 20/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.4770 - loss: 2.3188
Epoch 21/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.4925 - loss: 2.2187
Epoch 22/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.5114 - loss: 2.1239
Epoch 23/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.5257 - loss: 2.0466
Epoch 24/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.5382 - loss: 1.9759
Epoch 25/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.5516 - loss: 1.9074
Epoch 26/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.5618 - loss: 1.8515
Epoch 27/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.5751 - loss: 1.7843
Epoch 28/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.5834 - loss: 1.7408
Epoch 29/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.5942 - loss: 1.6854
Epoch 30/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.6022 - loss: 1.6497
Epoch 31/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.6122 - loss: 1.6019
Epoch 32/100			
3835/3835	<div></div>	84s	22ms/step - accuracy: 0.6188 - loss: 1.5692
Epoch 33/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6251 - loss: 1.5361
Epoch 34/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6307 - loss: 1.5145
Epoch 35/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6376 - loss: 1.4801
Epoch 36/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6398 - loss: 1.4636
Epoch 37/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6504 - loss: 1.4291
Epoch 38/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6529 - loss: 1.4041
Epoch 39/100			
3835/3835	<div></div>	85s	22ms/step - accuracy: 0.6552 - loss: 1.3991
Epoch 40/100			

3835/3835 ————— 85s 22ms/step - accuracy: 0.6603 - loss: 1.3692  
Epoch 41/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6647 - loss: 1.3548  
Epoch 42/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6697 - loss: 1.3377  
Epoch 43/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6710 - loss: 1.3197  
Epoch 44/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6739 - loss: 1.3091  
Epoch 45/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6807 - loss: 1.2867  
Epoch 46/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6815 - loss: 1.2789  
Epoch 47/100  
3835/3835 ————— 87s 23ms/step - accuracy: 0.6843 - loss: 1.2661  
Epoch 48/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6824 - loss: 1.2662  
Epoch 49/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6876 - loss: 1.2453  
Epoch 50/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6886 - loss: 1.2447  
Epoch 51/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6930 - loss: 1.2282  
Epoch 52/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.6937 - loss: 1.2273  
Epoch 53/100  
3835/3835 ————— 84s 22ms/step - accuracy: 0.6964 - loss: 1.2117  
Epoch 54/100  
3835/3835 ————— 84s 22ms/step - accuracy: 0.6995 - loss: 1.2020  
Epoch 55/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7003 - loss: 1.1905  
Epoch 56/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7018 - loss: 1.1909  
Epoch 57/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7035 - loss: 1.1760  
Epoch 58/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7062 - loss: 1.1693  
Epoch 59/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7079 - loss: 1.1629  
Epoch 60/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7101 - loss: 1.1594  
Epoch 61/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7091 - loss: 1.1481  
Epoch 62/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7126 - loss: 1.1439  
Epoch 63/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7160 - loss: 1.1315  
Epoch 64/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7142 - loss: 1.1342  
Epoch 65/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7177 - loss: 1.1269  
Epoch 66/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7161 - loss: 1.1251  
Epoch 67/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7180 - loss: 1.1172  
Epoch 68/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7198 - loss: 1.1156  
Epoch 69/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7179 - loss: 1.1121  
Epoch 70/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7214 - loss: 1.1072  
Epoch 71/100  
3835/3835 ————— 85s 22ms/step - accuracy: 0.7196 - loss: 1.1095  
Epoch 71: early stopping

```
In [33]: plt.plot(history.history["accuracy"],color="red",label="train")
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```



```
In [34]: plt.plot(history.history["loss"],color="red",label="train")
plt.title('loss function')
plt.ylabel('loss')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```



Predicting next word using this trained model.

```
In [36]: def predict_next(inp_text):
token_line=token.texts_to_sequences([inp_text])[0]
token_len=len(token_line)
if token_len<x_col:
    input=pad_sequences([token_line],maxlen=x_col, padding="pre")
elif token_len>x_col:
    input=token_line[-x_col:]
else:
    input=token_line
pred=model.predict(np.array([input]))
predict_index=np.argmax(pred)
for word,index in token.word_index.items():
    if index==predict_index:
```

```
print(word)
return
```

```
In [37]: predict_next("stood the statue of")
```

1/1  0s 282ms/step  
the

The original line was "...stood the statue of the Happy Prince."

```
In [39]: predict_next("when they finally plodded")
```

1/1  0s 31ms/step  
came

The original line was "...when they finally plodded home."

```
In [41]: predict_next("He is as beautiful")
```

1/1  0s 16ms/step  
as

The original line was "...He is as beautiful as a weathercock..."

So, this model is not predicting very well. The training accuracy is 71.96% which is not so good. However, we can use larger dataset and also more LSTM layers to get good accuracy and prediction if the system's capacity permits.